

Looking under the hood of the:

Parquet Format

André Kamman



@andrekamman

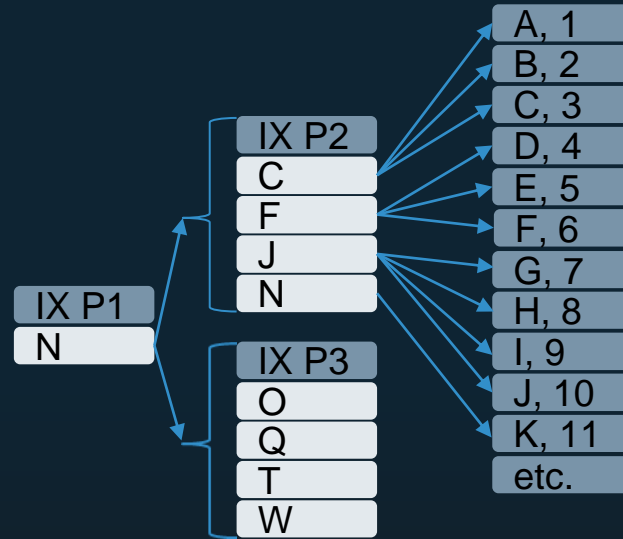


andrekamman@gmail.com

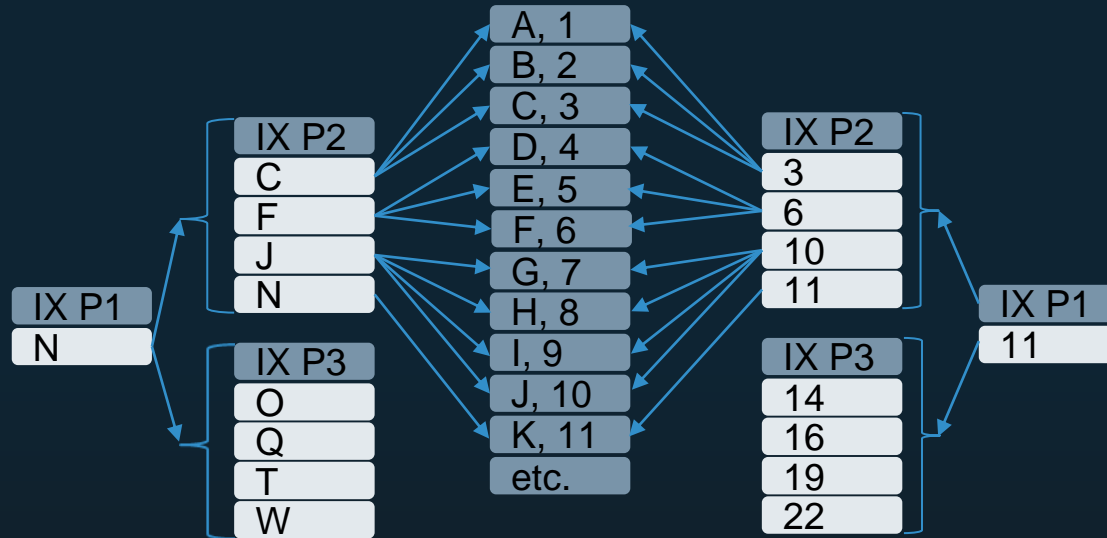
<https://evals.datagrillen.com/>




B-Tree



B-Tree



Spark + Databricks



bits

Python

Schedule ▾

democluster ▾

▾

▾

▾

Cmd 1

Python ▶ ▾

```
1 spark.read.format("csv") \  
2   .option("header", "true") \  
3   .option("inferSchema", "true") \  
4   .load("/mnt/lake/bits/videogames.csv") \  
5   .display()  
6
```

▶ (3) Spark Jobs

Table [Data Profile](#)

	Console ▲	Title
1	Nintendo DS	Super Mario 64 DS
2	Sony PSP	Lumines: Puzzle Fusion

```

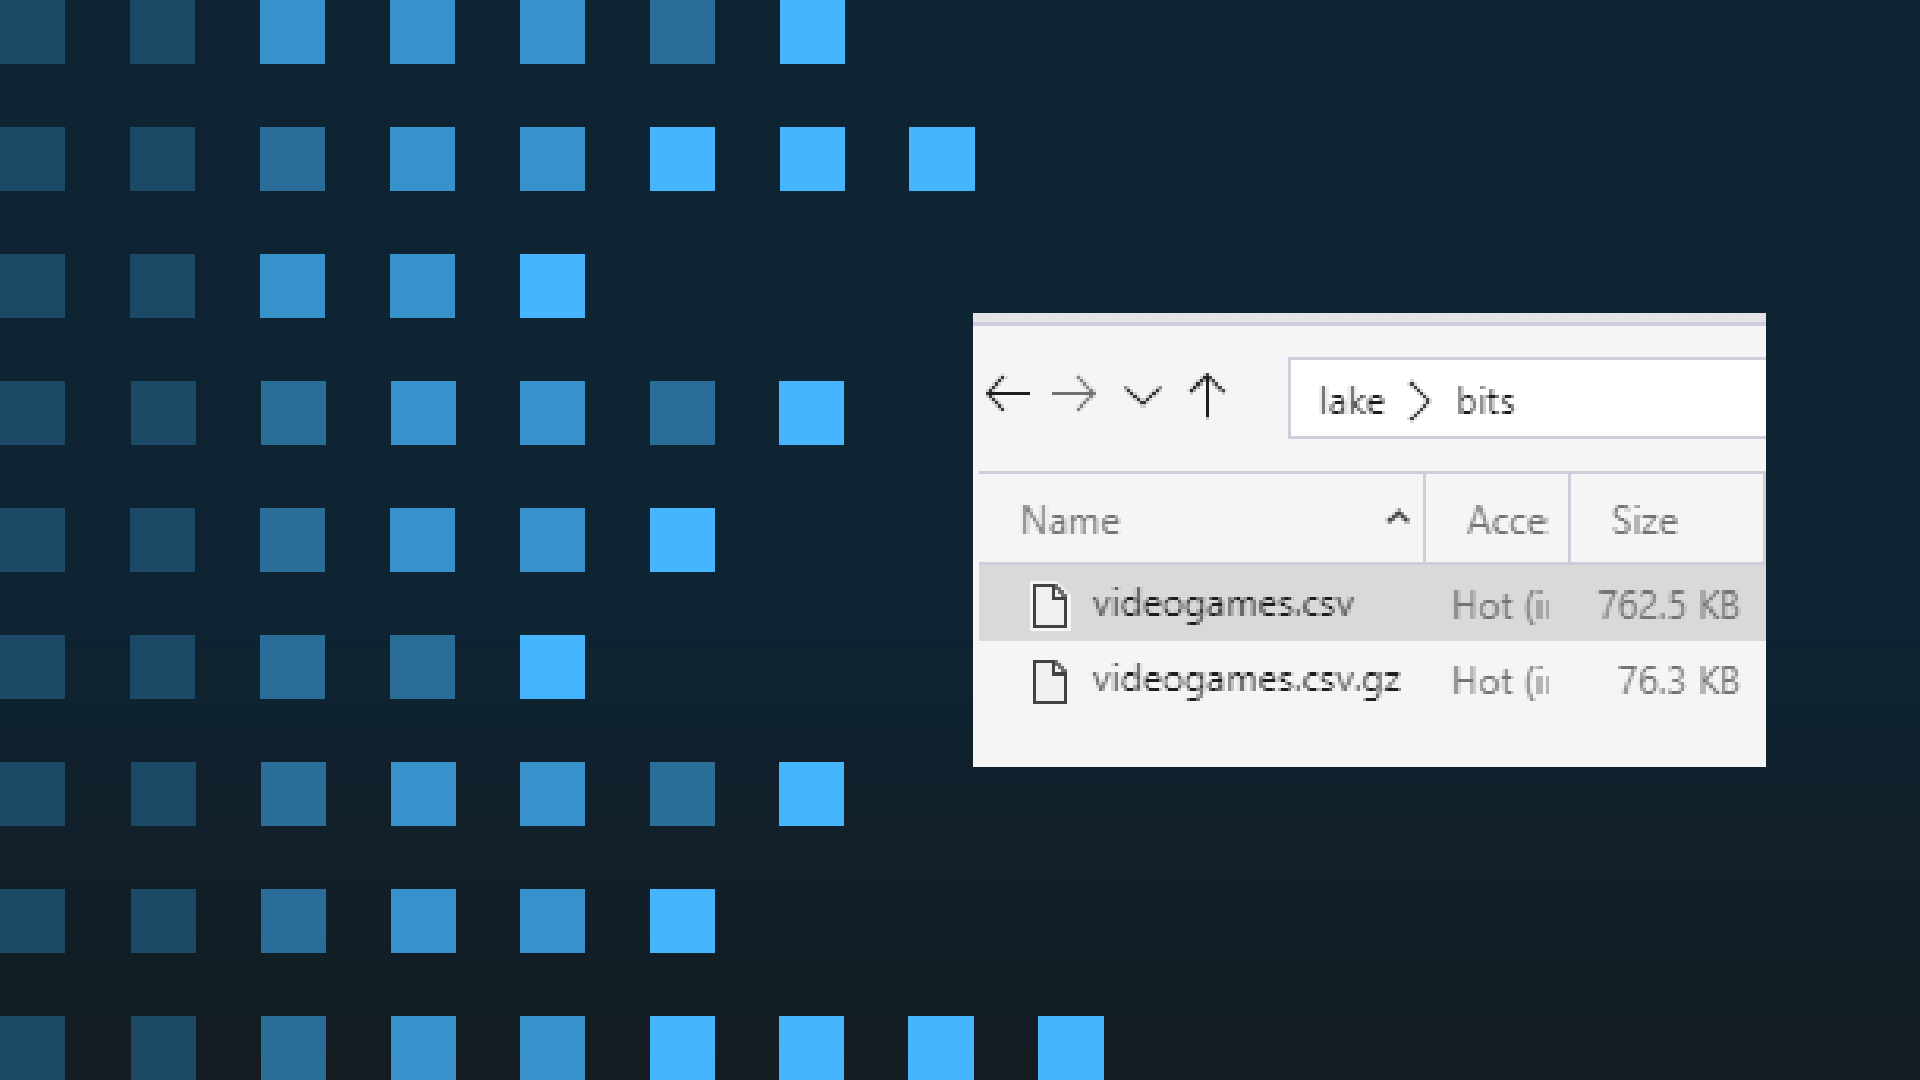
1 %sql
2 select
3     count(*) as Nr
4     ,Genre
5     ,Console
6 from videogames
7 group by Console, Genre
8 order by count(*) desc

```

► (2) Spark Jobs



Table [Data Profile](#)

	Nr ▲	Genre ▲	Console ▲
1	178	Action	X360
2	137	Action	PlayStation 3
3	123	Action	Nintendo DS
4	119	Action	Nintendo Wii
5	103	Action	Sony PSP
6	77	Sports	X360
7	58	Sports	PlayStation 3



Navigation icons: back, forward, up, down

lake > bits

Name	Acce	Size
 videogames.csv	Hot (ii	762.5 KB
 videogames.csv.gz	Hot (ii	76.3 KB

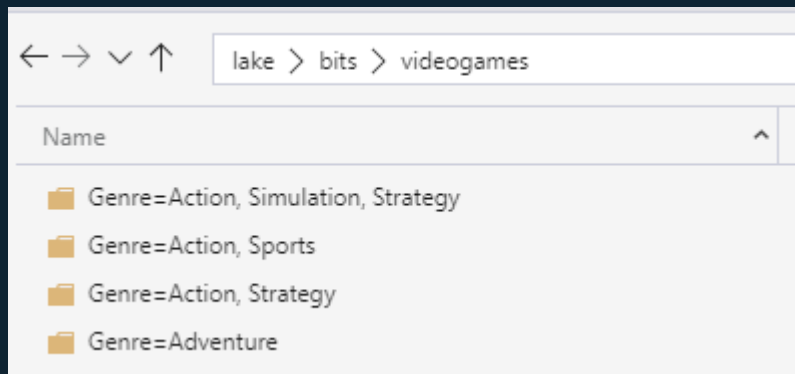
Possible Issues

- No schema (only column names)
- Have to uncompress on one node
- Have to read everything every time
- Only flat data, no nesting possible

Possible Issues

- No schema (only column names)
- Have to uncompress on one node
- **Have to read everything every time**
- Only flat data, no nesting possible

`.partitionBy("Genre")`



Finished calculating statistics for 'demolakestorage/lake/bits/videogames/'.
Active blob stats: 233 blobs, 142,384 bytes. Completed 07/03/2022 18:15.

```
1 %sql
2 select
3     count(*) as nr
4     ,Console
5 from videogames
6 where Genre = "Action"
7 group by Console
8 order by count(*) desc
```

▼ (2) Spark Jobs

- ▶ Job 46 [View](#) (Stages: 1/1)
- ▶ Job 47 [View](#) (Stages: 1/1, 1 skipped)

Table [Data Profile](#)

	nr	Console
1	178	X360
2	137	PlayStation 3
3	123	Nintendo DS
4	119	Nintendo Wii
5	103	Sony PSP

Submitted Time: 2022/03/07 17:18:54**Duration:** 0.5 s**Succeeded Jobs:** 46 47☐ Expand all the details in the query plan

Scan csv +details

Stages: 48.0

file sorting by size time	0 ms
number of files read	1
filesystem read data size	27.7 KiB
filesystem read data size (sampled)	55.4 KiB
filesystem read time (sampled)	26 ms
metadata time	2 ms
size of files read	27.7 KiB
number of partitions read	1
rows output	660

Possible Issues

- No schema (only column names)
- Have to uncompress on one node
- Have to read everything every time
- **Only flat data, no nesting possible**

A close-up of Jason Voorhees from the Friday the 13th movie series. He is wearing his signature white hockey mask with black ventilation holes and red markings around the eyes and mouth. He is dressed in a dark blue turtleneck and a brown, frayed jacket. He holds a large machete in his right hand, which is wearing a black glove. The background is a plain, light-colored wall. The word "JSON" is overlaid in the center in a bold, yellow, sans-serif font. Decorative blue squares of varying shades are arranged in a grid-like pattern in the top right, bottom left, and bottom right corners of the image.

JSON



Columns instead of Rows

	Col A	Col B	Col C	Col D
Row 0	A0	B0	C0	D0
Row 1	A1	B1	C1	D1
Row 2	A2	B2	C2	D2
Row 3	A3	B3	C3	D3
Row 4	A4	B4	C4	D4

	Col A	Col B	Col C	Col D
Row 0	A0	B0	C0	D0
Row 1	A1	B1	C1	D1
Row 2	A2	B2	C2	D2
Row 3	A3	B3	C3	D3
Row 4	A4	B4	C4	D4

A0	A1	A2	A3	A4	B1	B2	B3
B4	C0	C1	C2	C3	C4	D0	D1
D2	D3	D4					

	Col A	Col B	Col C	Col D
Row 0	A0	B0	C0	D0
Row 1	A1	B1	C1	D1
Row 2	A2	B2	C2	D2
Row 3	A3	B3	C3	D3
Row 4	A4	B4	C4	D4



Apache Parquet

Born out of a cooperation between Twitter & Cloudera
Open source storage format: PAX – Partition Attributes Across

Data Page Layouts for Relational Databases
on Deep Memory Hierarchies

Anastassia Ailamaki

David J. DeWitt

Mark D. Hill

The VLDB Journal 11 (2002)

PAR1

-

-

-

-

-

PAR1



PAR1

-

-

-

-

Footer Length (4 bytes)

PAR1





PAR1

-

-

-

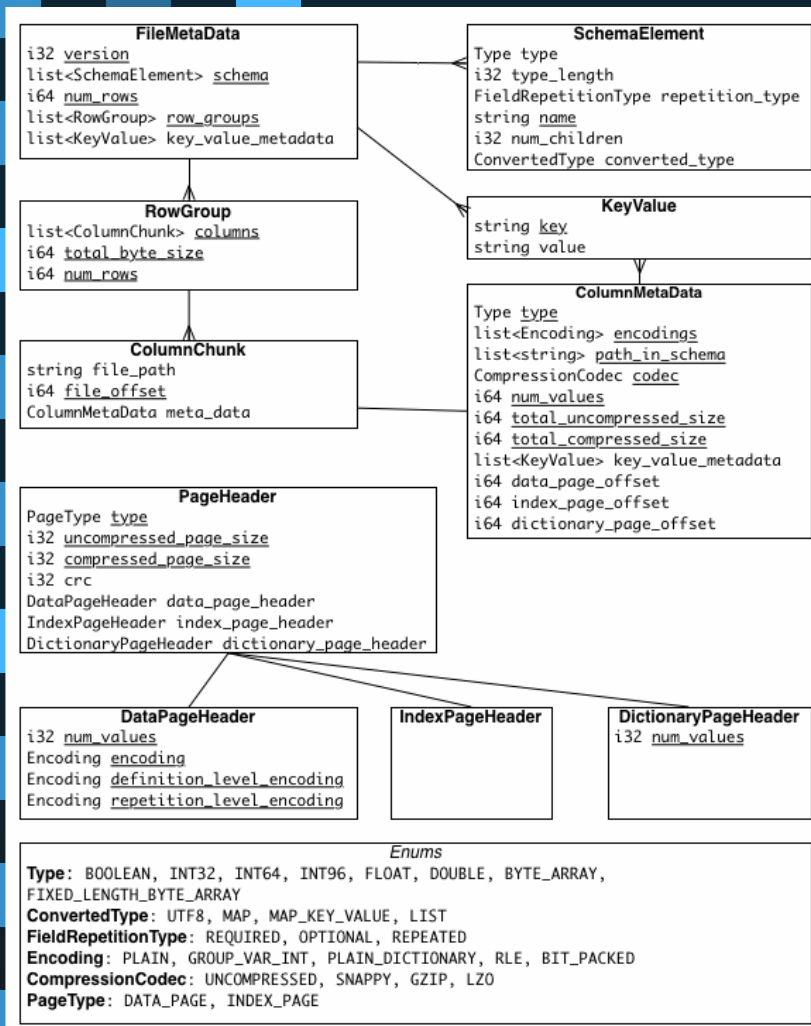
Footer (metadata)

Footer Length (4 bytes)

PAR1



Thrift Compact Protocol



PAR1

Row group 0

Column Chuck – Column A

Page 0 (Dictionary Page)

Page 1 (Data Page)

Chunk
–
Column
B

Chunk
–
Column
C

Chunk
–
Column
D

Row group 1

PAR1

Row group 0

Column Chuck – Column A

Page 0 (Dictionary Page)

Page 1 (Data Page)

Metadata

Min / Max / Count

Encoding and Compression Codec info

Repetition levels

(only used for nested columns)

Definition levels

(only used for nullable columns)

Encoded Data



3 Types of Metadata

- File metadata
- Column (chunk) metadata
- Page header metadata

2 Locations for metadata

Page Header
File Footer


Advantages So Far:

- Metadata at the end, one pass write.
- Column locations known, can grab only the columns needed.
- Min / Max values known, can skip pages, or even whole files.
- (Parquet table is usually a bunch of files and not just one)
- Directory based partitioning used to skip files, don't even have to parse the metadata.
- Data portion of pages are compressed, allowing compression without having to resort to a single node processor.

Encoding

Plain

- Fixed-width: back-to-back
- Variable Length: prefixed



RLE_Dictionary

- Run-length encoding
- Bit packing
- Dictionary Compression

Sony PSP
XBOX 360
SEGA
Nintendo Wii
Nintendo Wii
Nintendo Wii
Nintendo DS
Nintendo DS

0	Sony PSP
1	XBOX 360
2	SEGA
3	Nintendo Wii
4	Nintendo DS

0	1	2	3	3
3	4	4		

0	Sony PSP
1	XBOX 360
2	SEGA
3	Nintendo Wii
4	Nintendo DS


0	1	2	3,3	2,4
---	---	---	-----	-----

Delta_Length_Byte_Array



4,13,8,8,12,11

SegaPlaystation 3XBOX 360Sony PSPNintendo WiiNintendo DS



Delta_Binary_Packed




7	5	3	1	1	2	3	4
	-2	-2	-2	1	1	1	1
-2	0	0	0	3	3	3	3

Value Count : 8

First Value : 7

Minimum delta: -2

Bitwidth : 00 00 00 11 11 11



Delta_Binary_Packed

V2

1	2	3	4	5
	1	1	1	1
1	0	0	0	0


Value Count : 5

First Value : 1

Minimum delta: 1

Bitwidth : Not Needed!!

Adapted from “Decoding billions of integers per second through vectorization (D. Lemire, L.Boytsov)



Random Learnings

Dictionary pages are always PLAIN encoded

Delta Encoding with signed integers is called “ZigZag Encoding”

Row-Groups are 128MB, pages are 1MB by default (In spark I guess)

Is V2 The Future?

Maybe, a bit like IPv6 perhaps?
Check out Databricks Delta though!